

SENTENCE SELECTION FOR QUESTION ANSWERING

TUDOR BUZU, AI2



SENTENCE SELECTION FOR QA

- An important subtask of the question answering (QA) problem;
- The objective: identify the sentence from a large text or paragraphs, that contains the correct answer for a given question.

STRUCTURE OF THE PROJECT

- 1.) Structure of the dataset (analysis)**
- 2.) Data preprocessing**
 - 2.1) Context/passage sentence tokenization**
 - 2.2) Building labels for correct answers**
 - 2.3) Word tokenization**
 - 2.4) Lemmatization**
 - 2.5) Stopwords elimination**
- 3.) Solutions**
 - 3.1) Random approach**
 - 3.2) String kernels**
 - 3.2.1) Spectrum kernel**
 - 3.2.2) Presence kernel**
 - 3.2.3) Intersection kernel**
 - 3.3) Sentence similarity**
 - 3.4) Logistic regression using pretrained word embeddings**
 - 3.5) BM25**
 - 3.6) Combined methods**
- 4.) Results**

STRUCTURE OF THE DATASET (ANALYSIS)

1.) SQuAD (Stanford Question Answering Dataset)[1]

- a new reading comprehension dataset that covers a wide range of topics (from musical celebrities, sports, history to abstract concepts);
- 536 articles (sampled uniformly at random from the top 10,000 articles of English Wikipedia);
- the articles were partitioned randomly into training (80%), development (10%) and test (10%) sets.
- 23,215 paragraphs;
- up to 5 questions on the content of one paragraph;
- the correct answer is a segment of text from the corresponding reading passage.

DATA PREPROCESSING

- Context/passage sentence tokenization;
- Building labels for correct answers;
- Word tokenization;
- Lemmatization;
- Stopwords elimination.

SOLUTIONS

- Random approach;
- String kernels;
 - Spectrum kernel;
 - Presence kernel;
 - Intersection kernel;
- Sentence similarity;
- Logistic regression using pretrained word embeddings;
- BM25;
- Combined methods;

RANDOM APPROACH

- A random permutation for candidate answers is generated;
- This method gives weak results: $\text{Prec@1} \sim 0.25$ and $\text{MAP} \sim 0.5$.

SPECTRUM KERNEL

- This kernel is defined like the sum of products between frequencies of question and candidate answer common words;
- Results: Prec@1 ~ 0.765 and MAP ~ 0.86.

```
def spectrum_kernel_value(question_words, sentence_words):  
    kernel_value = 0  
    vocab_inters = set(question_words).intersection(sentence_words)  
  
    for word in vocab_inters:  
        kernel_value += num(word, question_words) * num(word, sentence_words)  
  
    return kernel_value
```


PRESENCE KERNEL

- This kernel is defined like the number of common words in question and candidate answer;
- Results: Prec@1 ~ 0.8 and MAP ~ 0.87.

```
def presence_kernel_value(question_words, sentence_words):  
    kernel_value = 0  
    vocab_inters = set(question_words).intersection(sentence_words)  
  
    return len(vocab_inters)
```

INTERSECTION KERNEL

- This kernel is defined like the sum of minimum frequencies of common words in question and candidate answer;
- Results: Prec@1 ~ 0.8 and MAP ~ 0.87.

```
def intersection_kernel_value(question_words, sentence_words):  
    kernel_value = 0  
    vocab_inters = set(question_words).intersection(sentence_words)  
  
    for word in vocab_inters:  
        kernel_value += min(num(word, question_words), num(word, sentence_words))  
  
    return kernel_value
```

SENTENCE SIMILARITY

- Similarity between a question and a candidate answer is calculated as the average of maximum similarities between synonyms of each word from question and synonyms from answer;
- Results: Prec@1 ~ 0.74 and MAP ~ 0.83.

RESULTS

	Train set		Dev set	
Method \ Evaluation metric	Prec@1	MAP	Prec@1	MAP
Random approach	0.245	0.491	0.261	0.500
Spectrum kernel	0.757	0.856	0.779	0.866
Presence kernel	0.7853	0.8717	0.8180	0.8866
Intersection kernel	0.7854	0.8718	0.8181	0.8865
Sentence similarity (path_similarity)	0.717	0.826	0.750	0.841
Sentence similarity (wup_similarity)	0.707	0.819	0.741	0.835

LOGISTIC REGRESSION

- I trained a logistic regression classifier using word embeddings from a sentence;
- I used pretrained word embeddings (Google News Dataset);
- A sentence is the average of its word embeddings;
- Feature vector for training: diff_abs, diff_sqr, dot_product, min, max, sum of question and candidate answer sentence embeddings.

RESULTS

	Train set		Dev set	
Feature vector \ Evaluation metric	Prec@1	MAP	Prec@1	MAP
diff_abs	0.730	0.835	0.752	0.8444
diff_sqr	0.727	0.833	0.752	0.8442
dot_product	0.655	0.793	0.677	0.802
min	0.714	0.825	0.739	0.836
max	0.713	0.825	0.739	0.836
sum	0.280	0.524	0.294	0.535

BM25

- A good ranking function used by search engines to rank matching documents according to their relevance to a given search query.
- Given a query Q , containing keywords q_1, q_2, \dots, q_n , the BM25 score of a document D is:

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)},$$

Usually, the values of free parameters: $k_1 \in [1.2, 2.0]$ and $b = 0.75$.

- Results:
 - ❖ train: Prec@1 ~ 0.777 and MAP ~ 0.869.
 - ❖ dev: Prec@1 ~ 0.808 and MAP ~ 0.882.

COMBINED METHODS

- I used the scores from the previous methods and trained again a logistic regression classifier;
- Scores from string kernels + LR (word embeddings) + BM25;
- The results look better:
 - ❖ train: Prec@1 ~ 0.804 and MAP ~ 0.883.
 - ❖ dev: Prec@1 ~ 0.838 and MAP ~ 0.898.

FINAL RESULTS

	Train set		Dev set	
Method \ Evaluation metric	Prec@1	MAP	Prec@1	MAP
Random approach	0.245	0.491	0.261	0.500
Spectrum kernel	0.757	0.856	0.779	0.866
Presence kernel	0.7853	0.8717	0.8180	0.8866
Intersection kernel	0.7854	0.8718	0.8181	0.8865
Sentence similarity (path_similarity)	0.717	0.826	0.750	0.841
Sentence similarity (wup_similarity)	0.707	0.819	0.741	0.835
LR (word embeddings,feature vector = dif_abs)	0.730	0.835	0.752	0.8444
BM25	0.777	0.869	0.808	0.882
LR(LR(word embeddings) + string kernels + BM25)	0.804	0.883	0.838	0.898

Q&A

